

Dev Soni

New Delhi, India | +91-7398399009 | devsagar12380@gmail.com |
[linkedin.com/in/dev-soni-8b664625a](https://www.linkedin.com/in/dev-soni-8b664625a) | github.com/Dev-2004-DA

M.Sc. Statistics student (CGPA 8.36, University of Delhi) with hands-on expertise in Python, MySQL, and end-to-end ML pipelines across finance, healthcare, and operations. Adept at translating statistical models into actionable business insights and deploying live applications. Seeking a data analyst / data science internship where statistical rigour drives measurable decisions.

Academic Qualifications

Passing Year	Course	Institute	Percentage/CGPA
2027	M.Sc. Statistics	Ramjas College, University of Delhi	8.36
2025	B.Sc. (Hons) Statistics	Institute of Science, Banaras Hindu University	7.95
2022	Class XII, UP Board	Government Inter College, Sultanpur	73.8%
2020	Class X, UP Board	Government Inter College, Sultanpur	79%

Core Competencies

Time Series Forecasting | ML Pipelines | Statistical Inference & Hypothesis Testing | Feature Engineering | EDA & Data Cleaning | Regression & Classification | Model Evaluation (ROC-AUC, CV) | SQL & ETL Design | Data Visualisation | Streamlit Deployment | Predictive Modelling | Business Insight Communication

Academic & Industry Projects

Intel Image Classification — CNN to Transfer Learning | TensorFlow · Keras · Keras Tuner · TensorBoard · Transfer Learning Jun 2026

Built and compared three CNN architectures for 6-class natural scene classification, progressing from a custom CNN to Transfer Learning with full error analysis.

- Designed a custom Sequential CNN (Conv2D → BatchNorm → ReLU → MaxPooling → SpatialDropout, filters 16→128) and trained on the **Intel Image Classification dataset** (~14K train / ~3K test / ~7.3K unlabelled images, 6 scene classes).
- Ran two independent **Keras Tuner** hyperparameter search runs (with and without L2 regularization); both converged to the same **77.8% test accuracy** ceiling, diagnosing the architecture as the bottleneck and motivating Transfer Learning.
- Implemented progressive fine-tuning of **MobileNetV2** across 3 runs with ReduceLROnPlateau-driven LR scheduling; best run (last 20 layers unfrozen) achieved **90.57% test accuracy** (88.37% train), with validation consistently exceeding training accuracy.
- Diagnosed an input-resolution mismatch with EfficientNet (B0/B3/V2B3) at 150x150 (45-50% accuracy); switched to **Xception**, which converged in 9 epochs to **91.73% test accuracy** (93.55% train, 1.82% gap).
- Used **TensorBoard** throughout to monitor accuracy/loss curves in real time, diagnose over/underfitting, and identify optimal ReduceLROnPlateau trigger points across all three models.
- Built confusion matrices and classification reports for all 3 models; identified a consistent **Glacier↔Mountain** misclassification pattern across every architecture, driven by genuine visual similarity.
- Diagnosed and fixed a critical **tf.data.Dataset shuffling bug** that misaligned labels and predictions across separate iterations (producing a misleading 17% accuracy); fixed by collecting both within a single pass.
- Ran all 3 models on 7,301 unlabelled images: **75% unanimous agreement** across architecturally distinct models; conducted a confidence-calibration analysis on disagreement cases, finding MobileNetV2 was often overconfident (up to 100%) on incorrect predictions vs. Xception's better-calibrated errors.
- Result: 77.8% (scratch CNN) → 90.57% (MobileNetV2, ~14MB) → 91.73% (Xception, ~112MB) — a **13.93% improvement**, with MobileNetV2 offering near-equivalent accuracy at ~8x smaller model size.

→ github.com/Dev-2004-DA/Intel-Image-Classification

Google Stock Price Forecasting | SARIMAX · Random Forest · XGBoost · Python

Feb – Mar 2026

Engineered a signal-generation system to automate directional entry-point detection from 2 years of OHLCV data.

- Built an end-to-end **time series forecasting pipeline** using Google (GOOGL) stock data, covering EDA → stationarity testing → feature engineering → statistical & ML modelling.
- Conducted EDA on 2 years of daily OHLCV data (**500+ trading sessions**); performed **Augmented Dickey-Fuller (ADF) test** to identify stationarity and determine optimal differencing ($d = 1$).
- Applied **Granger causality** analysis to select significant exogenous variables and engineered optimal lag features (**High_lag6, Low_lag9, Open_lag7**).
- Designed and executed **parallelised SARIMAX grid search**, identifying optimal model: **SARIMAX(0,1,0)(2,1,0,20)**; achieved **RMSE = 22.87, MAPE = 6.70%**.
- Diagnosed ML model failure ($R^2 = -10.13$) due to **regime shift (train-test distribution mismatch)**, demonstrating strong debugging ability.
- Recovered performance by **engineering return-based features** (percentage change) to handle non-stationarity, including rolling volatility, intraday & gap returns, and multi-period momentum; boosted to **$R^2 = 0.84+$** .
- Implemented **TimeSeriesSplit cross-validation** to prevent data leakage; trained and tuned Random Forest, Gradient Boosting, XGBoost, Voting Regressor; best: **Random Forest Test $R^2 = 84.89\%$** .
- Key insight: **feature engineering (returns) > model tuning** in non-stationary financial time series.

→ github.com/Dev-2004-DA/google-stock-forecasting | Live App: t-i-m-e-series-forecasting.streamlit.app

Cardiovascular Disease Prediction | Scikit-learn · Random Forest · ROC-AUC · Python

Jan – Feb 2026

Developed an early-warning classifier across 70,000 patient records to support clinical screening decisions.

- Performed domain-driven data cleaning on **70,000+ patient records**, correcting physiological inconsistencies in blood pressure and BMI values using clinical constraints.
- Applied stratified train-test split and **5-fold cross-validation** to ensure unbiased model evaluation.
- Benchmarked **10+ classifiers**; selected Random Forest achieving **~75–76% CV accuracy, OOB ~73%**. Identified overfitting via train vs. CV gaps; diagnosed bias-variance trade-offs across models.
- Feature importance analysis revealed **systolic/diastolic BP as dominant predictors (~70% combined)**; documented SVM scalability constraints influencing model selection.
- Performed ROC analysis and threshold tuning to boost **disease recall from 69% → 85%** for screening-oriented deployment; analysed precision-recall trade-offs to align with clinical risk tolerance.

→ github.com/Dev-2004-DA/cardiovascular-disease-prediction | Live App: cardiovascular-diseases-prediction.streamlit.app

Medical Insurance Cost Prediction | Scikit-learn · ElasticNet · Random Forest · Python

Jan – Feb 2026

Built end-to-end regression ML pipelines on 1,338 patient records for insurance pricing support.

- Performed extensive EDA, feature engineering, outlier analysis, and categorical encoding using **Pipelines and ColumnTransformers**.
- Discovered that the target variable (**charges**) is **bimodal (Smokers vs. Non-Smokers)**, explaining non-linearity and residual non-normality in linear models.
- Applied **GridSearchCV + 5-fold CV** to tune hyperparameters across Linear Regression, ElasticNet, Polynomial, Decision Tree, and Random Forest models.
- Achieved **~86% test R^2** with Random Forest; minimal train-test gap proving robustness on high-variance data.
- Conducted residual diagnostics and feature importance analysis; established **smoking status as primary cost driver (~70% feature importance)**, supporting data-driven pricing decisions.

→ github.com/Dev-2004-DA/medical-insurance-cost-prediction

House Price Prediction | Scikit-learn · Ridge/Lasso · PCA · Python

Jan – Feb 2026

Built regression ML pipelines on a 21K-row housing dataset.

- Applied feature engineering, **regularisation (Ridge/Lasso)**, Polynomial features, and **PCA** to manage multicollinearity and bias-variance tradeoff.
- Evaluated models using train-test split, **5-fold cross-validation**, and residual diagnostics; achieved **~69% CV R^2 ($\pm 1.3\%$)**, demonstrating stable generalisation beyond the linear baseline.

→ github.com/Dev-2004-DA/ML_beginning

Airbnb NYC Data Cleaning Pipeline | MySQL · Window Functions · ETL

Sep – Oct 2025

Designed a production-grade SQL ETL pipeline for large-scale data standardisation.

- Designed a **3-stage ETL pipeline (raw → staging → clean)** on **100K+ listings**; applied ROW_NUMBER, COALESCE, TRIM, STR_TO_DATE for deduplication & standardisation across multiple tables.
- Applied advanced SQL Window & String Functions for missing value imputation and data quality improvement.

→ github.com/Dev-2004-DA/airbnb-data-cleaning-sql

Café Sales Data Cleaning | MySQL · Window Functions · ETL

Oct – Nov 2025

Cleaned and standardised a transactional dataset for operational analysis.

- Cleaned a café sales dataset of **10K+ rows** using MySQL, handling duplicates, missing values, and inconsistent entries.
- Standardised columns, enforced constraints, and corrected data types; used SQL window functions (LAG, COALESCE) to impute missing values.

→ github.com/Dev-2004-DA/cafe-sales-data-cleaning-sql

Synthetic Customer Profiling | Python · Pandas · Matplotlib · Seaborn

Jun 2025

Exploratory data analysis and customer segmentation on synthetic retail data.

- Used Pandas, NumPy for data cleaning & analysis; Matplotlib and Seaborn for comprehensive visualisation.

→ github.com/Dev-2004-DA

Tools & Skills

Languages & Tools: Python (NumPy, Pandas, Matplotlib, Seaborn, Scipy, Statsmodels, Scikit-learn, TensorFlow, Keras), MySQL, MS Excel (Pivot Tables, Power Query, XLOOKUP, Basic Dashboards), Power BI, Git/GitHub, Streamlit

Machine Learning: Supervised Learning (Linear/Logistic Regression, Random Forest, Gradient Boosting, XGBoost, SVM), Ridge/Lasso/ElasticNet, Model Evaluation (CV, ROC-AUC, Precision-Recall), Feature Engineering, ANN, Bias-Variance Tradeoff

Statistics & Forecasting: Time Series (ARIMAX, SARIMAX, VARMAX), Hypothesis Testing, ADF Test, Granger Causality, Residual Diagnostics, Statistical Inference, Sampling Theory

Deep Learning: ANN, CNN (TensorFlow, Keras), Transfer Learning (MobileNetV2, Xception, EfficientNet), Keras Tuner, TensorBoard, Bias-Variance Tradeoff, Model Evaluation

Educational Background

M.Sc. Statistics, Ramjas College, University of Delhi | CGPA: 8.36

2025 – 2027 (ongoing)

Probability Theory · Time Series · Statistical Inference · Design of Experiments · Stochastic Processes · Non-Parametric Inference · Statistical Computing (R & Python)

B.Sc. (Hons) Statistics, Institute of Science, Banaras Hindu University, Varanasi | CGPA: 7.95

2022 – 2025

Statistical Methods · Statistical Inference · Sample Surveys · Design of Experiments · Decision Theory · Operations Research · Programming with C

Class XII, UP Board, Government Inter College, Sultanpur

2022

Percentage: 73.8%

Class X, UP Board, Government Inter College, Sultanpur

2020

Percentage: 79%

Certifications & Activities

- **Deloitte Australia Data Analytics Job Simulation** · Forage, Aug 2025 — simulated end-to-end client data analysis & dashboard reporting engagement.
- Presented research on **Stochastic Modelling of Reliance Industries stock prices** at academic seminar, Ramjas College — 5 models from Random Walk to GBM.
- **Activities:** Badminton · Swimming